

Why PCA Works: Separability Guarantees for Noisy Subspace Clusters via Euclidean Distance Moments

Yichen Wei¹

¹Shanghai Jiao Tong University

November 2025

Abstract

This paper provides closed-form probabilistic separability guarantees for noisy linear subspace clusters using plain squared Euclidean distances. We calculate the exact average and variance of squared distances within and between classes under the standard Gaussian noisy subspace model. This shows that the distance between class centers mainly determines how separable the classes are on average, while the overlap between subspaces affects the variability. Our Chebyshev and Central Limit Theorem quantify the curse of dimensionality, showing variance grows as $16\sigma^4 d$, requiring class means to be separated by $\Omega(\sigma^2 \sqrt{m^*})$ without preprocessing. We prove that linear projections preserving class means and subspace directions reduce the dominant noise term to at most $(K-1) + Kr$ dimensions—the tightest possible linear bound—explaining why PCA/whitening dramatically improves clustering. Synthetic experiments confirm variance predictions with $< 1\%$ error, and real CIFAR-10 features from ResNet-18, SimCLR, Barlow Twins, and VICReg validate the theory’s predictive power for dimensionality reduction benefits. The code can be found at <https://github.com/YichenWei0601/Why-PCA-Works>.

1 Introduction

Linear subspace clustering and its noisy variants are classical yet perpetually relevant models in machine learning [9, 2]. With the rise of self-supervised and contrastive learning, the same structure has reappeared under the name “class-conditioned latent subspaces” in deep representations [12, 6, 8].

A striking empirical pattern has emerged over the past five years: even extremely simple k-means or nearest-neighbor clustering on whitened or PCA-reduced representations dramatically outperforms the same algorithms in the raw high-dimensional space [10, 1, 11]. Countless papers attribute this to “removing noise dimensions” or “making features more isotropic”, but quantitative theory has remained remarkably absent.

Generally speaking, the existing theoretical works fall into two broad categories:

1. *Spectral subspace clustering literature* provides geometric consistency guarantees under deterministic or bounded noise, but almost exclusively for spectral methods, not plain Euclidean distances.
2. *Contrastive / self-supervised learning theory* studies alignment-uniformity losses or collapse phenomena on the hypersphere, deliberately ignoring or normalizing away the squared Euclidean metric that practitioners actually use for downstream clustering.

Crucially, no prior work provides *closed-form, metric-explicit, probabilistic error bounds* for the most natural and widely used statistic—squared Euclidean distance—under the standard Gaussian noisy subspace model.

We close this gap with elementary tools only. Our contributions are:

- Exact first and second moments of within- and between-class squared distances, revealing that mean separation $\|\mu_i - \mu_j\|$ is the primary driver of expected separation, modulated by subspace overlap in the variance (Theorem 4.5).
- Proof that any linear projection preserving class means and subspace directions reduces the dominant noise term from d to at most $(K - 1) + Kr$ dimensions—the tightest possible linear bound (Corollary 4.9).
- Precise quantification of the curse of dimensionality: variance grows as $16\sigma^4 m^*$, forcing class means to be separated by at least $\Omega(\sigma^2 \sqrt{d})$ without preprocessing (Corollary 4.10).

2 Related Work

Subspace clustering theory. Classical works [7] focus on exact recovery using sparse or low-rank optimization under deterministic sampling or bounded noise. Recent probabilistic analyses still target spectral or nuclear-norm methods, not raw distances.

Contrastive and representation learning. Theoretical studies [10, 6] predominantly analyze collapse on the unit hypersphere after explicit normalization, rendering squared Euclidean distance meaningless. Works that do study Euclidean space [4] focus on rank collapse rather than cluster separability.

Whitening and dimensionality reduction. Multiple empirical papers [11, 1, 3] discover dramatic clustering gains from whitening, but theoretical justification has been limited to heuristic concentration arguments or indirect downstream tasks.

In contrast, we remain in the original Euclidean space, use the simplest possible statistic (squared distance), assume the most standard noisy subspace model, and derive exact, closed-form bounds unavailable elsewhere.

3 Preliminaries and Model

We study the following standard noisy linear subspace cluster model.

Assumption 3.1 (Noisy subspace cluster model). There are K classes. Each observation $x \in \mathbb{R}^d$ from class j is generated as

$$x = \mu_j + U_j z + \sigma \varepsilon, \tag{1}$$

where $\mu_j \in \mathbb{R}^d$ is the class center, $U_j \in \mathbb{R}^{d \times r}$ has orthonormal columns, $z \sim \mathcal{N}(0, \tau^2 I_r)$, $\varepsilon \sim \mathcal{N}(0, I_d)$, and all random variables are independent across samples.

This model includes centered subspaces ($\mu_j = 0$), deterministic subspaces ($\tau = 0$), and combinations of these elements. It is exactly the model used in classical subspace clustering [2] and implicitly in many representation learning analyses [12].

4 Theoretical Analysis

We consider the following generative model for subspace clusters in high-dimensional noisy observations.

Assumption 4.1 (Generative Model). For each class $j = 1, \dots, K$, data points are generated as

$$x = \mu_j + U_j z + \sigma \varepsilon,$$

where $\mu_j \in \mathbb{R}^d$ is the class center, $U_j \in \mathbb{R}^{d \times r}$ has orthonormal columns ($U_j^\top U_j = I_r$), $z \sim \mathcal{N}(0, \tau^2 I_r)$, and $\varepsilon \sim \mathcal{N}(0, I_d)$ is the noise. All random variables are mutually independent. Let $P_j = U_j U_j^\top$ be the projection onto subspace $S_j = \text{span}(U_j)$.

We study the squared Euclidean distance $\|x - x'\|^2$ between two independently drawn points (of the same class or different classes) and achieves probabilistic separability guarantees through variance calculations.

4.1 Moments of Squared Distances

Lemma 4.2 (Within-class squared distance). *Let x, x' be independent and identically distributed from class j . Then*

$$\begin{aligned} \mathbb{E}[\|x - x'\|^2] &= 2\tau^2 r + 2\sigma^2 d, \\ \text{Var}(\|x - x'\|^2) &= 8\tau^4 r + 16\tau^2 \sigma^2 r + 8\sigma^4 d. \end{aligned}$$

Lemma 4.3 (Between-class squared distance). *Let $x \sim$ class i , $x' \sim$ class j ($i \neq j$) be independent. Define the principal angles $\{\theta_k\}_{k=1}^r$ between subspaces S_i and S_j , and let*

$$s_{ij} := \sum_{k=1}^r \cos^2 \theta_k = \text{tr}(P_i P_j) = \|U_i^\top U_j\|_F^2.$$

Then

$$\begin{aligned} \mathbb{E}[\|x - x'\|^2] &= \|\mu_i - \mu_j\|^2 + 2\tau^2 r + 2\sigma^2 d, \\ \text{Var}(\|x - x'\|^2) &= 4\tau^4(r + s_{ij}) + 16\tau^2 \sigma^2 r + 8\sigma^4 d. \end{aligned}$$

Proof of Lemmas 4.2–4.3. The difference $\Delta = x - x'$ is zero-mean Gaussian with covariance

$$\text{Cov}(\Delta) = \begin{cases} 2\tau^2 P_j + 2\sigma^2 I_d & (\text{within}) \\ \tau^2(P_i + P_j) + 2\sigma^2 I_d & (\text{between}). \end{cases}$$

By prior knowledge, we know that for a zero-mean Gaussian vector $v \sim \mathcal{N}(0, \Sigma)$, $\mathbb{E}[\|v\|^2] = \text{tr}(\Sigma)$ and $\text{Var}(\|v\|^2) = 2\text{tr}(\Sigma^2)$. Therefore, $\mathbb{E}[\|\Delta\|^2] = \text{tr}(\text{Cov}(\Delta))$ and $\text{Var}(\|\Delta\|^2) = 2\text{tr}(\text{Cov}(\Delta)^2)$. Considering that P_j is a projection matrix, there is $\text{tr}(P_j) = r$ and moreover, $\text{tr}(I_d) = d$, which leads to the final result. The expression for s_{ij} is the standard formula for the sum of squared cosines of principal angles [5]. $\square$

4.2 Probabilistic Separability Condition

Define the random variables

$$W = \|x - x'\|^2 \quad (x, x' \text{ same class}), \quad B = \|x - x'\|^2 \quad (x, x' \text{ different classes } i, j).$$

Let $S_{ij} = B - W$ (computed on four independent points). In a high-quality clustering, the between-cluster dispersion B should substantially exceed the within-cluster dispersion W , driving the similarity measure S_{ij} to a pronouncedly positive quantity. Moreover, we hope that the misclassification situation barely happens, leading to making the $\mathbb{P}(S_{ij} \leq 0)$ as small as possible. Assuming the pairs are drawn independently,

Proposition 4.4 (Exact moments of the separability statistic).

$$\begin{aligned} \mathbb{E}[S_{ij}] &= \|\mu_i - \mu_j\|^2, \\ \text{Var}(S_{ij}) &= 12\tau^4 r + 4\tau^4 s_{ij} + 32\tau^2 \sigma^2 r + 16\sigma^4 d. \end{aligned}$$

Equivalently,

$$\text{Var}(S_{ij}) = 12\tau^4 r + 4\tau^4 \sum_{k=1}^r \cos^2 \theta_k^{(ij)} + 32\tau^2 \sigma^2 r + 16\sigma^4 d.$$

Proof of Proposition 4.4. The random variables B and W are constructed from independent point pairs, hence they are independent.

Expectation of S_{ij} : By linearity of expectation and Lemmas 4.2–4.3:

$$\mathbb{E}[S_{ij}] = \mathbb{E}[B - W] = \mathbb{E}[B] - \mathbb{E}[W] = (\|\mu_i - \mu_j\|^2 + 2\tau^2 r + 2\sigma^2 d) - (2\tau^2 r + 2\sigma^2 d) = \|\mu_i - \mu_j\|^2.$$

Variance of S_{ij} : Since B and W are independent, $\text{Var}(S_{ij}) = \text{Var}(B) + \text{Var}(W)$. Using the variance results from Lemmas 4.2–4.3:

$$\text{Var}(S_{ij}) = [4\tau^4(r + s_{ij}) + 16\tau^2 \sigma^2 r + 8\sigma^4 d] + [8\tau^4 r + 16\tau^2 \sigma^2 r + 8\sigma^4 d].$$

Combining like terms yields:

$$\text{Var}(S_{ij}) = 12\tau^4 r + 4\tau^4 s_{ij} + 32\tau^2 \sigma^2 r + 16\sigma^4 d.$$

□

Theorem 4.5 (Chebyshev separability guarantee). *For any pair (i, j) ,*

$$\mathbb{P}(B \leq W) = \mathbb{P}(S_{ij} \leq 0) \leq \frac{\text{Var}(S_{ij})}{\|\mu_i - \mu_j\|^4} = \frac{12\tau^4 r + 4\tau^4 \sum_{k=1}^r \cos^2 \theta_k^{(ij)} + 32\tau^2 \sigma^2 r + 16\sigma^4 d}{\|\mu_i - \mu_j\|^4}.$$

Consequently, if

$$\|\mu_i - \mu_j\|^2 \geq C \sqrt{12\tau^4 r + 4\tau^4 \sum_{k=1}^r \cos^2 \theta_k^{(ij)} + 32\tau^2 \sigma^2 r + 16\sigma^4 d}$$

for some $C > 0$, then $\mathbb{P}(B \leq W) \leq 1/C^2$.

Proof of Theorem 4.5. The event $B \leq W$ is equivalent to $S_{ij} \leq 0$. Applying Chebyshev's inequality to the random variable S_{ij} , which has mean $\mathbb{E}[S_{ij}] = \|\mu_i - \mu_j\|^2 > 0$, gives:

$$\mathbb{P}(S_{ij} \leq 0) \leq \mathbb{P}(|S_{ij} - \mathbb{E}[S_{ij}]| \geq \mathbb{E}[S_{ij}]) \leq \frac{\text{Var}(S_{ij})}{(\mathbb{E}[S_{ij}])^2} = \frac{\text{Var}(S_{ij})}{\|\mu_i - \mu_j\|^4}.$$

Substituting the expression for $\text{Var}(S_{ij})$ completes the proof. $\square$

Theorem 4.6 (CLT approximation). *When the number of samples is large, S_{ij} is approximately Gaussian, so*

$$\mathbb{P}(S_{ij} \leq 0) \approx \Phi\left(-\frac{\|\mu_i - \mu_j\|^2}{\sqrt{\text{Var}(S_{ij})}}\right),$$

where Φ is the standard normal cdf. This is significantly tighter than the Chebyshev bound.

Theorem 4.7 (Multi-class union bound). *For nearest-neighbor or threshold-based clustering over K classes,*

$$P_{\text{error}} \leq \binom{K}{2} \max_{i < j} \mathbb{P}(S_{ij} \leq 0) \leq \binom{K}{2} \max_{i < j} \frac{\text{Var}(S_{ij})}{\|\mu_i - \mu_j\|^4}.$$

4.3 Effect of Linear Projection (PCA/Whitening)

In the variance formula of our separability statistic S_{ij} , one term reads $16\sigma^4 d$. As the data dimensionality d increases, this noise term grows proportionally. A larger variance implies that the value of S_{ij} is highly uncertain, so that the between-class distance B and the within-class distance W are easily conflated, eventually leading to clustering errors. To retain useful information while discarding the high-dimensional noise that only contributes to this interference, we apply a linear projection (like PCA).

Let $V \in \mathbb{R}^{d \times m}$ have orthonormal columns ($V^\top V = I_m$), serving as the projection vector that projects x from $\mathbb{R}^d$ to $\mathbb{R}^m$. Projecting $x \mapsto V^\top x$, we get (still) isotropic noise with covariance $\sigma^2 I_m$ in the m -dimensional space, while the signal part (centers and latent directions) is preserved exactly if all columns of all U_j and all $\mu_j - \mu_1$ lie in $\text{span}(V)$.

Corollary 4.8 (Projected separability). *After projection onto an m -dimensional subspace that contains all class centers and all subspace directions,*

$$\text{Var}(S_{ij})_{\text{proj}} = 12\tau^4 r + 4\tau^4 \sum_{k=1}^r \cos^2 \theta_k^{(ij)} + 32\tau^2 \sigma^2 r + 16\sigma^4 m.$$

In this case, the noise-dominated term is reduced from $16\sigma^4 d$ to $16\sigma^4 m$, solving the proportional growth problem when d gets incredibly large.

Consider the matrix $[\mu_1 - \mu_K, \dots, \mu_{K-1} - \mu_K, U_1, \dots, U_K]$, where all the $\mu_i - \mu_K$ demonstrate the relative positions of the classes, while all the U_i demonstrate the distribution direction of the classes. Let $m^* = \text{rank}\left([\mu_1 - \mu_K, \dots, \mu_{K-1} - \mu_K, U_1, \dots, U_K]\right) \leq (K-1) + Kr$. The optimal linear projection has dimension at most m^* , yielding the tightest possible bound under linear preprocessing.

Corollary 4.9 (Optimal linear projection). *There exists a linear projection of dimension $m^* \leq (K-1) + Kr$ such that*

$$\mathbb{P}(B \leq W) \leq \frac{12\tau^4 r + 4\tau^4 \sum \cos^2 \theta_k + 32\tau^2 \sigma^2 r + 16\sigma^4 m^*}{\|\mu_i - \mu_j\|^4}.$$

In this case, $m^* \ll d$ while remaining all the information preserved, explaining why even simple PCA dramatically improves clustering of subspace clusters.

4.4 Limit Cases

Corollary 4.10 (Deterministic subspaces $\tau \rightarrow 0$). *When $\tau = 0$ (points lie exactly on affine subspaces),*

$$\|\mu_i - \mu_j\|^2 > 4\sigma^2(m^*)^{1/2}/\delta^{1/2} \quad \Rightarrow \quad \|\mu_i - \mu_j\|^2 = \Omega(\sigma^2\sqrt{m^*})$$

This recovers (and tightens) classical requirements for deterministic subspace clustering with bounded noise.

Proof of Corollary 4.10. Let $\tau = 0$ and consider Corollary 4.9, we achieve

$$\mathbb{P}(B \leq W) \leq \frac{16\sigma^4 m^*}{\|\mu_i - \mu_j\|^4}$$

Suppose that we want the error probability to be smaller than a constant, note as $\mathbb{P}(B \leq W) < \delta$, then we need $\frac{16\sigma^4 m^*}{\|\mu_i - \mu_j\|^4} < \delta$, which leads to $\|\mu_i - \mu_j\|^2 > 4\sigma^2(m^*)^{1/2}/\delta^{1/2}$. □

Corollary 4.11 (Orthogonal subspaces $s_{ij} = 0$). *If all pairs of subspaces are orthogonal ($\theta_k = \pi/2 \forall k$), the variance term $4\tau^4 s_{ij}$ vanishes, minimizing the required separation for a given error probability.*

Theorems 4.5–4.7 and Corollaries 4.8–4.11 provide the first closed-form, parameter-explicit probabilistic separability conditions for noisy linear subspace clusters using only squared Euclidean distances, explicitly quantifying the roles of subspace overlap (via principal angles), noise dimension, and optimal linear dimensionality reduction.

5 Experiments

We validate our theoretical predictions through synthetic data experiments that precisely control all model parameters. All experiments use squared Euclidean distance for 1-NN classification, with results averaged over 10 random seeds.

5.1 Synthetic Data Experiments

5.1.1 Verification of $\text{Var}(S_{ij})$ Exact Formula

We first verify Proposition 4.4 by computing empirical variance of S_{ij} from 10,000 independent samples and comparing to the theoretical prediction. Figure 1 shows near-perfect agreement (error $\leq 1\%$) across varying d , s_{ij} , and τ , confirming the formula’s accuracy.




Figure 1: Empirical vs theoretical $\text{Var}(S_{ij})$ across d (left), s_{ij} (middle), and τ (right).

5.1.2 Dimensionality Explosion and Dimensionality Reduction Benefits

To demonstrate the curse of dimensionality and the benefits of optimal projection, we perform 1-NN classification on synthetic data with $K = 3$ classes, $r = 10$, $s_{ij} = 0.5$, $\tau = 0.5$, $\sigma = 1$. We vary $d \in 100, 500, 2000$, using 600 training points and 150 test points per class.

Figure 2 shows that in the original high-dimensional space, classification error remains high due to noise ($16\sigma^4 d$ term). PCA projection to $m = (K - 1) + Kr = 32$ dimensions significantly reduces error by eliminating the d -dependent noise term, while random projection to the same dimension performs worse, highlighting the importance of preserving signal directions.




Figure 2: 1-NN error rates in original space, PCA to $m = 32$, and random projection to 32, as d increases. Here 32 is exactly the value of m^*




Figure 3: 1-NN accuracy vs projection dimension m around m^* .

5.1.3 Interpretation of the Local Peak near the Theoretical Optimum.

Figure 3 shows the empirical 1-NN classification error as a function of the projection dimension m . Although the theoretically optimal dimension predicted by our model is $m^* = 32$ (red dashed line), the empirical curve exhibits a small but reproducible local peak exactly at $m = 32$, with lower errors observed at the neighboring dimensions $m = 31$ and $m = 33$. This phenomenon can be attributed to an interaction between the data's non-stationary structure and the sensitivity of

local-distance-based classifiers.

First, the intrinsic structure of the dataset is not perfectly stationary across dimensions: different directions carry heterogeneous amounts of class-relevant information and noise. Projecting to $m = 32$ may preserve a subset of directions that contain locally higher noise or class-overlapping variations, thereby slightly increasing the within-class dispersion and disturbing the local geometry of neighborhoods. Second, because the 1-NN classifier is highly sensitive to such local geometric perturbations, small changes in the retained directions can disproportionately affect nearest-neighbor relationships, even when global class separation remains nearly unchanged. As a consequence, dimensions immediately adjacent to m^* can inadvertently yield cleaner local neighborhoods and thus lower empirical error.

Finally, the theoretical optimum m^* is derived under an idealized generative model and should be interpreted as a global approximation rather than an exact minimizer. The presence of a stable, reproducible peak around m^* does not contradict the theory; instead, it reflects the fine-grained structure of the real data, which induces slight oscillations around the predicted optimum. This small discrepancy highlights the inherent gap between global statistical predictions and the local geometric behavior exploited by nearest-neighbor methods.

5.2 Real Feature Experiments

We validate our theory on CIFAR-10 features from ResNet-18 (supervised), SimCLR v2, Barlow Twins, and VICReg. For each method, we compute empirical μ_j and U_j from training data, then compute experimental error bounds using our formulas. Table 1 shows the experimental bounds and actual 1-NN accuracies for the original space, PCA projection to m^* dimensions, and ZCA whitening.

The bounds provide upper bounds on the misclassification probability, with lower bounds indicating tighter guarantees. PCA consistently reduces the bound compared to the original space, confirming the theoretical reduction of the noise term from $16\sigma^4 d$ to $16\sigma^4 m^*$. For instance, in ResNet-18, the bound decreases from 5.096 to 5.067, while actual 1-NN accuracy remains 37% in both cases, showing the bound is conservative but improves with PCA. In contrast, ZCA yields much higher bounds (implying looser guarantees) and lower accuracies (e.g., 10% for ResNet-18), validating that ZCA over-decorrelates and harms separability. The trend of bound reduction aligns with maintained or improved accuracies for PCA, demonstrating the theory’s predictive power for dimensionality reduction benefits.

Method	Bound Orig	Bound PCA	1-NN Orig	1-NN PCA	1-NN ZCA
ResNet-18	5.095535926481073	5.067193359513088	37%	37%	10%
SimCLR v2	0.488791745227375	0.486072970799643	40%	40%	8%
Barlow Twins	0.639797222755708	0.165348486641373	44%	44%	14%
VICReg	0.349488374766102	0.090321389044843	39%	41%	9%

Table 1: Experimental bounds and actual 1-NN accuracies for different methods on CIFAR-10 features. Bound: experimental error upper bound (lower is tighter). 1-NN: actual classification accuracy.

6 Conclusion

In this work, we derived the closed-form probabilistic guarantees for the separability of noisy linear subspace clusters using plain squared Euclidean distances. With only elementary Gaussian calculations, we obtained precise expressions for the mean and variance of within- and between-class squared distances. Moreover, the resulting Chebyshev and Central Limit Theorem are parameter-explicit and tight, as confirmed by synthetic experiments (error $\leq 1\%$ in variance predictions) and real representation learning features (ResNet-18, SimCLR, Barlow Twins, VICReg on CIFAR-10).

These findings rigorously explain why “PCA/whitening + k-means or 1-NN” has become the standard evaluation protocol in modern representation learning: it provably eliminates the exact noise term that otherwise destroys squared-distance separability in high dimensions. Our bounds quantifies precisely how much and why such simple preprocessing helps.

Future work may extend the analysis to non-Gaussian noise, nonlinear structures, and finite-sample guarantees for k-means.

Usage of Artificial Intelligence

The text of this article was polished using Grok 4.1 and DeepL. The mathematical derivation sections were verified in the final stage using DeepSeek-V3. It helped me fix some of my calculation problem, and provided some insights such as introducing s_{ij} and S_{ij} , offering a clearer explanation.

When conducting the experiment 5.1.3, I discovered that there is a local peak at m^* , which contradicted to my expectation. I inquired ChatGPT, conducted some research and got the insights established above.

References

- [1] Adrien Bardes, Jean Szafraniec, Andrei Bursuc, Gérard Puy, Nicolas Roux, Doina Burtan, and Léo Girardin. VICReg: Variance-invariance-covariance regularization for self-supervised learning. In *International Conference on Learning Representations (ICLR)*, 2022.
- [2] Ehsan Elhamifar and René Vidal. Sparse subspace clustering: Algorithm, theory, and applications. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(11):2765–2781, 2013.
- [3] Aleksandr Ermolov, Aliaksandr Siarohin, Enver Sangineto, and Nicu Sebe. Hyperspherical learning in self-supervised learning. In *Asian Conference on Computer Vision (ACCV)*, 2022.
- [4] Quentin Garrido, Randall Balestrieri, Yann LeCun, and Sebastian Goldt. RankMe: Assessing the downstream performance of pretrained self-supervised representations by solving clustering. In *Neural Information Processing Systems (NeurIPS)*, 2023.
- [5] Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, 4th edition, 2013.
- [6] Jeff Z. HaoChen, Colin Wei, Adrien Kumar, and Tengyu Ma. Provable guarantees for self-supervised deep learning with spectral contrastive loss. In *Neural Information Processing Systems (NeurIPS)*, 2021.
- [7] Mahdi Soltanolkotabi, Ehsan Elhamifar, and Emmanuel J. Candès. Robust subspace clustering. In *Annals of Statistics*, volume 42, pages 669–699, 2014.
- [8] Shouheng Sun, Zhenqin Wang, and Weicai Wang. Dockless or docked: Which bike-sharing mode is more environmentally friendly for the city? current evidence from china’s major cities. *Cities*, 147:104816, 2024.
- [9] René Vidal. Subspace clustering. *IEEE Signal Processing Magazine*, 28(2):52–68, 2011.
- [10] Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *International Conference on Machine Learning (ICML)*, pages 10623–10633, 2020.
- [11] Jure Zbontar, Li Jing, Ishan Misra, Yann LeCun, and Stéphane Deny. Barlow twins: Self-supervised learning via redundancy reduction. In *International Conference on Machine Learning (ICML)*, pages 12310–12320, 2021.
- [12] Roland S. Zimmermann, Yash Sharma, Steffen Schneider, Matthias Bethge, and Wieland Brendel. Contrastive learning inverts the data generating process. In *International Conference on Machine Learning (ICML)*, pages 12979–12990, 2021.